



Crate lancedb

Search Summary

LanceDB is an open-source database for vector-search built with persistent storage, which greatly simplifies retrieval, filtering and management of embeddings.

The key features of LanceDB include:

- Production-scale vector search with no servers to manage.
- Store, query and filter vectors, metadata and multi-modal data (text, images, videos, point clouds, and more).
- Support for vector similarity search, full-text search and SQL.
- Native Rust, Python, Javascript/TypeScript support.
- Zero-copy, automatic versioning, manage versions of your data without needing extra infrastructure.
- GPU support in building vector indices¹.
- Ecosystem integrations with LangChain 🦜🔗, LlamaIndex 🐾, Apache-Arrow, Pandas, Polars, DuckDB and more on the way.

Getting Started

LanceDB runs in process, to use it in your Rust project, put the following in your `Cargo.toml`:

`cargo add lancedb`



Crate Features

Experimental Features

These features are not enabled by default. They are experimental or in-development features that are not yet ready to be released.

- `remote` - Enable remote client to connect to LanceDB cloud. This is not yet fully implemented and should not be enabled.

Quick Start

Connect to a database.

```
let db = lancedb::connect("data/sample-lancedb").execute().await.unwrap();
```

LanceDB accepts the different form of database path:

- /path/to/database - local database on file system.
- s3://bucket/path/to/database or gs://bucket/path/to/database - database on cloud object store
- db://dbname - Lance Cloud

You can also use [ConnectOptions] to configure the connection to the database.

```
use object_store::aws::AwsCredential;
let db = lancedb::connect("data/sample-lancedb")
    .aws_creds(AwsCredential {
        key_id: "some_key".to_string(),
        secret_key: "some_secret".to_string(),
        token: None,
    })
    .execute()      :::
    .await
    .unwrap();
```

LanceDB uses [arrow-rs](#) to define schema, data types and array itself. It treats `FixedSizeList<Float16/Float32>` columns as vector columns.

For more details, please refer to the [LanceDB documentation](#).

Create a table

To create a Table, you need to provide a `arrow_schema::Schema` and a `arrow_array::RecordBatch` stream.

```
use arrow_array::{RecordBatch, RecordBatchIterator};
use arrow_schema::{DataType, Field, Schema};

let schema = Arc::new(Schema::new(vec![
    Field::new("id", DataType::Int32, false),
    Field::new(
```

```

"vector",
    DataType::FixedSizeList(Arc::new(Field::new("item", DataType::Float,
        true,
    )),
])
// Create a RecordBatch stream.
let batches = RecordBatchIterator::new(
    vec![RecordBatch::try_new(
        schema.clone(),
        vec![
            Arc::new(Int32Array::from_iter_values(0..256)),
            Arc::new(
                FixedSizeListArray::from_iter_primitive::<Float32Type, _,
                (0..256).map(|_| Some(vec![Some(1.0); 128])),
                128,
            ),
        ),
    ],
)
.unwrap()]
.into_iter()
.map(Ok),
    schema.clone(),
);
db.create_table("my_table", Box::new(batches))
.execute()
.await
.unwrap();

```

Create vector index (IVF_PQ)

LanceDB is capable to automatically create appropriate indices based on the data types of the columns. For example,

- If a column has a data type of `FixedSizeList<Float16/Float32>`, LanceDB will create a IVF-PQ vector index with default parameters.

```
use lancedb::index::Index;
tbl.create_index(&["vector"], Index::Auto)
    .execute()
    .await
    .unwrap();
```

User can also specify the index type explicitly, see [Table::create_index](#).

Open table and search

```
let results = table
    .query()
    .nearest_to(&[1.0; 128])
    .unwrap()
    .execute()
    .await
    .unwrap()
    .try_collect::<Vec<_>>()
    .await
    .unwrap();
```

1. Only in Python SDK. ↩

Re-exports

```
pub use connection::ConnectNamespaceBuilder;
pub use connection::Connection;
pub use error::Error;
pub use error::Result;
pub use table::Table;
pub use connection::connect;
pub use connection::connect_namespace;
```

Modules

[arrow](#)

[connection](#)

[data](#)

Functions to establish a connection to a LanceDB database

Data types, schema coercion, and data cleaning and etc.

[database](#)
[dataloader](#)
[embeddings](#)
[error](#)
[index](#)
[io](#)
[ipc](#)
[query](#)
[rerankers](#)
[table](#)
[utils](#)

The database module defines the `Database` trait and related types.

IPC support

LanceDB Table APIs

Structs

[ObjectStoreRegistry](#)

A registry of object store providers.

[Session](#)

Re-export Lance Session and ObjectStoreRegistry for custom session creation A user session holds the runtime state for a [crate::Dataset](#)

Enums

[DistanceType](#)

